



MAJOR: DATA SCIENCE (Computational Track)

Project: Transfer Learning for Custom Datasets in the Small-Data Regime

Project Number: 20(Basement)

Project Member 1

Name : Harshitha Saripilli
NJIT UCID : HS759
E-MAIL : hs59@njit.edu

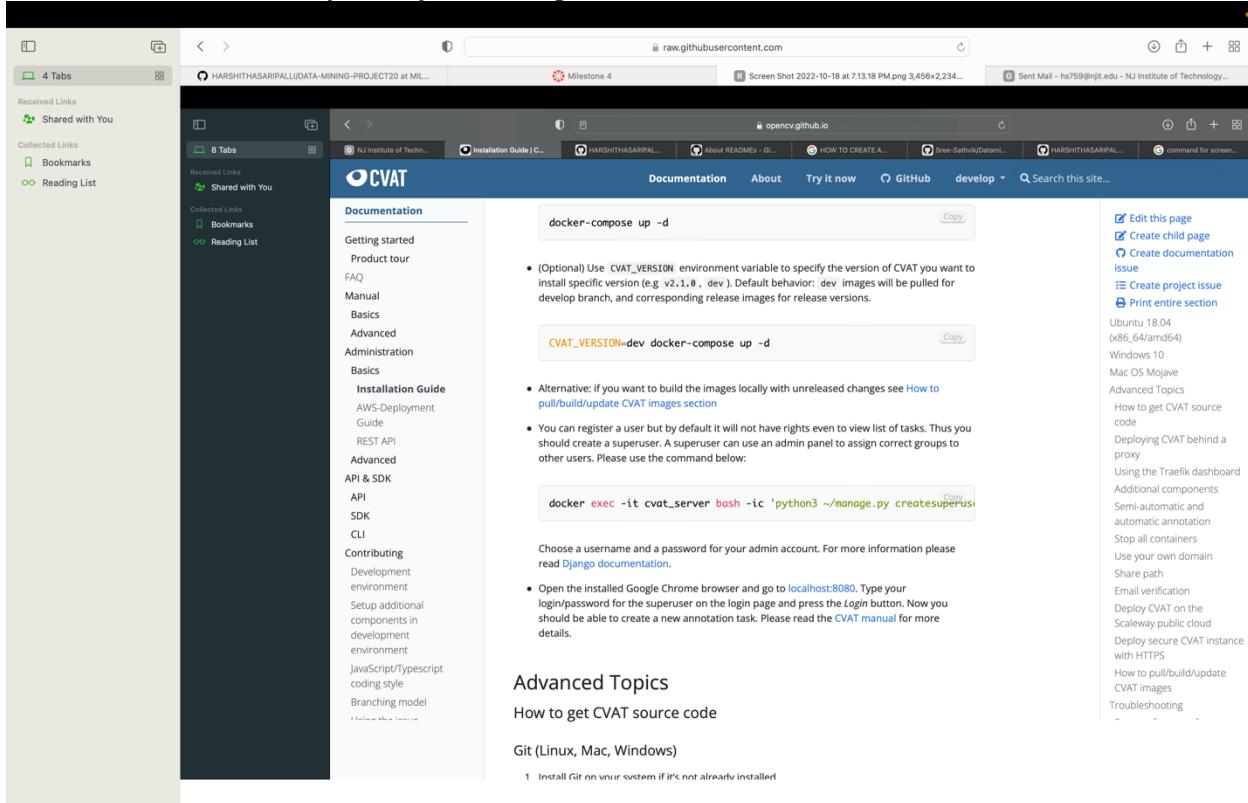
Project Member 2

Name : Sree Sathvik Chilamkurthy
NJIT UCID : SC2493
E-MAIL : sc2493@njit.edu

Professor : Pantelis Monogioudis

MILESTONE 1

In this stage you will be working with CVAT - you need to install CVAT locally to your laptop/desktop and we urge you to follow the docker setup instructions [here](#) PS: You can also create an account on the [cvat.ai](#) site but we don't offer any guarantees that the project can be delivered in its entirety with a cvat instance that is not managed by you. Submit a GitHub repository with a branch titled 'milestone-1' with the readme file containing the CVAT installation instructions you followed and a screenshot of your computer of the login screen. Add as collaborator the TA.



The screenshot shows the Docker Desktop interface with a tutorial window open. The title is "First, clone a repository". It explains that the Getting Started project is a simple GitHub repository containing everything needed to build an image and run it in a container. It provides two command examples:

```
docker run --name repo alpine/git clone https://github.com/docker/getting-started.git >> docker cp repo:/git/getting-started/;
```

You can also type the command directly in a command line interface.

At the bottom, there are "Next Step" and "Skip tutorial" buttons.

The screenshot shows the CVAT Documentation page for Mac OS Mojave. The left sidebar lists various documentation sections: Getting started, Product tour, FAQ, Manual, Basics, Advanced, Administration, and Contributing. The main content area is titled "Mac OS Mojave" and contains instructions for installing Git on Mac OS. It includes a command example:

```
git --version
```

If you don't have it installed already, it will prompt you to install it. More instructions can be found [here](#).

- Download Docker for Mac. Double-click Docker.dmg to open the installer, then drag Moby the whale to the Applications folder. Double-click Docker.app in the Applications folder to start Docker. More instructions can be found [here](#).
- There are several ways to install Git on a Mac. The easiest is probably to install the Xcode Command Line Tools. On Mavericks (10.9) or above you can do this simply by trying to run git from the Terminal the very first time.

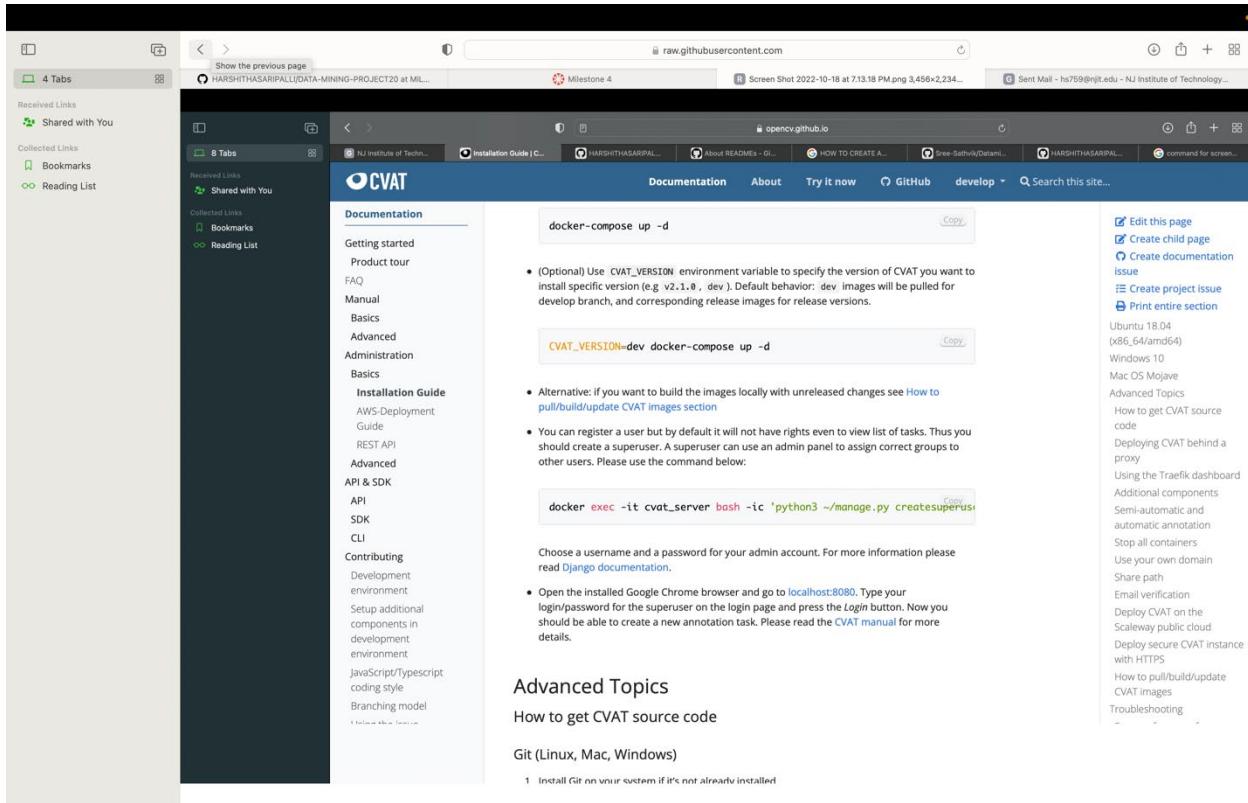
The following command will clone the latest develop branch:

```
git clone https://github.com/opencv/cvat  
cd cvat
```

See [alternatives](#) if you want to download one of the release versions or use the `wget` or `curl` tools.

- Run docker containers. It will take some time to download the latest CVAT release and other required images like postgres, redis, etc. from DockerHub and create containers.

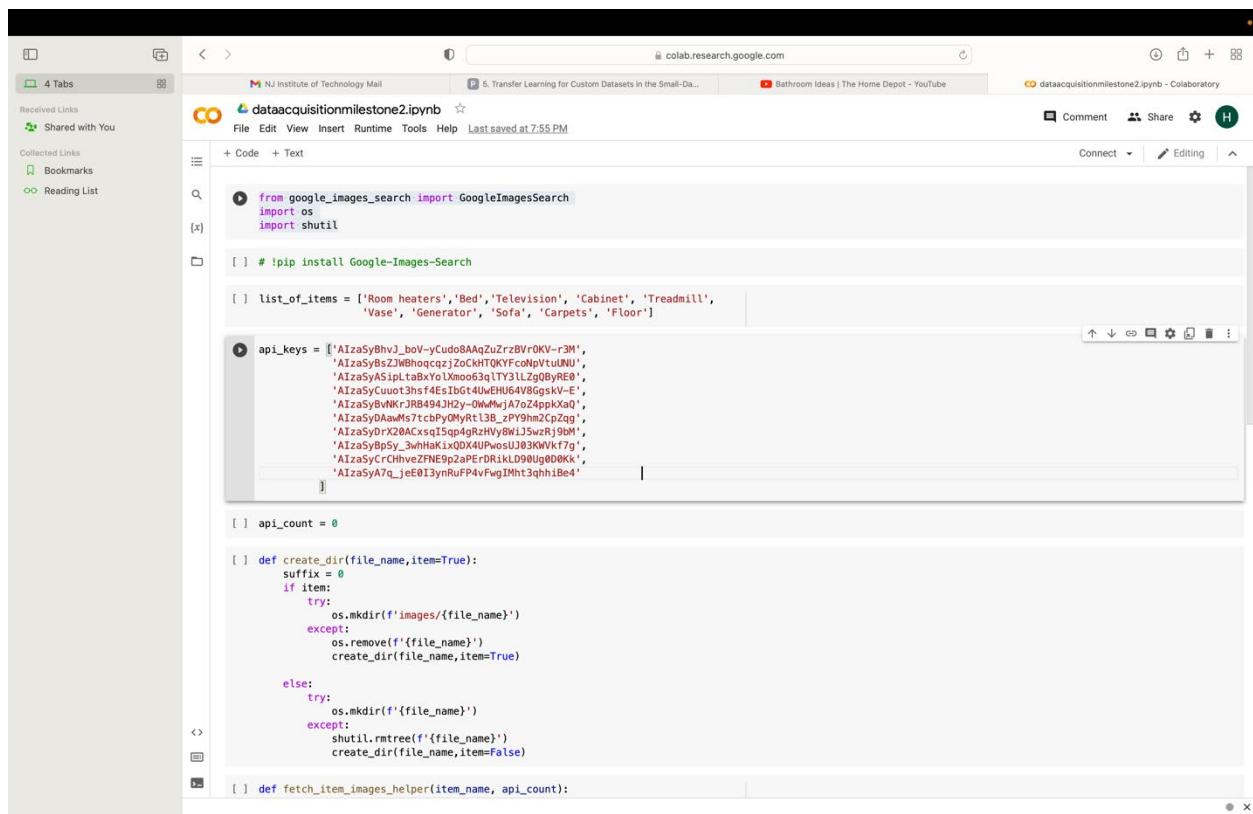
The right sidebar contains links for editing the page, creating child pages, documentation issues, project issues, and various system configurations.



MILESTONE-2

5.2. Milestone 2: Data Acquisition (20 points)

Project Team 1 is assigned home area 1 (Kitchen) just like Project Team 6, Project Team 3 is assigned home area 3 (Bedroom) just like Project Team 8 and so on ...Select 10 product categories (eg sinks, faucets) that HD sells in each room and that are shown on their promotional video. You will need to scrape using [beautiful soup](#) or retrieve using the [Google Custom Search API](#)(better solution), 100 images of each of the selected product categories. Make sure you select products that are not included on the video, meaning if a specific sink is shown on the video, you scrape / search different but similar sinks. Submit a Github branch titled 'data-acquisition' and a markdown file called [data-acquisition.md](#), containing the code that achieve the milestone and a pytest test file that can show that the code works.



The screenshot shows a Google Colaboratory notebook titled "dataacquisitionmilestone2.ipynb". The code in the notebook is as follows:

```
from google_images_search import GoogleImagesSearch
import os
import shutil

# !pip install Google-Images-Search

list_of_items = ['Room heaters', 'Bed', 'Television', 'Cabinet', 'Treadmill',
                 'Vase', 'Generator', 'Sofa', 'Carpets', 'Floor']

api_keys = ['AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M',
            'AIzaSyhj_1_pvY-CudoAAAq2UzrzBvR0Kv-r3M']

api_count = 0

def create_dir(file_name, item=True):
    suffix = 0
    if item:
        try:
            os.mkdir(f'images/{file_name}')
        except:
            os.remove(f'{file_name}')
            create_dir(file_name, item=True)
    else:
        try:
            os.mkdir(f'{file_name}')
        except:
            shutil.rmtree(f'{file_name}')
            create_dir(file_name, item=False)

def fetch_item_images_helper(item_name, api_count):
```

The screenshot shows a Google Colab notebook titled "dataacquisitionmilestone2.ipynb". The code in the cell is as follows:

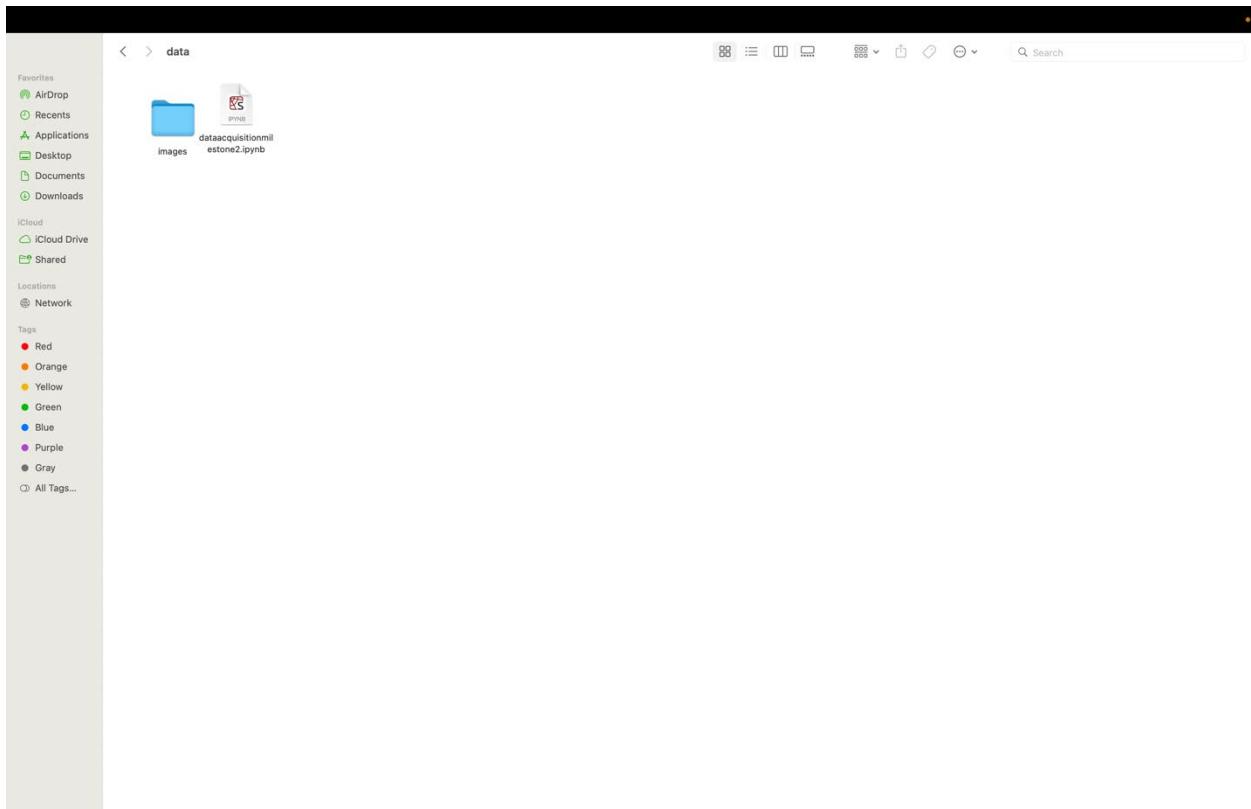
```
[ ] def fetch_item_images_helper(item_name, api_count):
    _search_params = {
        'q': f'{item_name}',
        'num': 100, 'fileType': 'jpg|png',
    }
    gis = GoogleImagesSearch(api_keys[api_count], 'b2dad1940a4a74e80')
    gis.search(search_params=_search_params, path_to_dir=f'images/{item_name}')

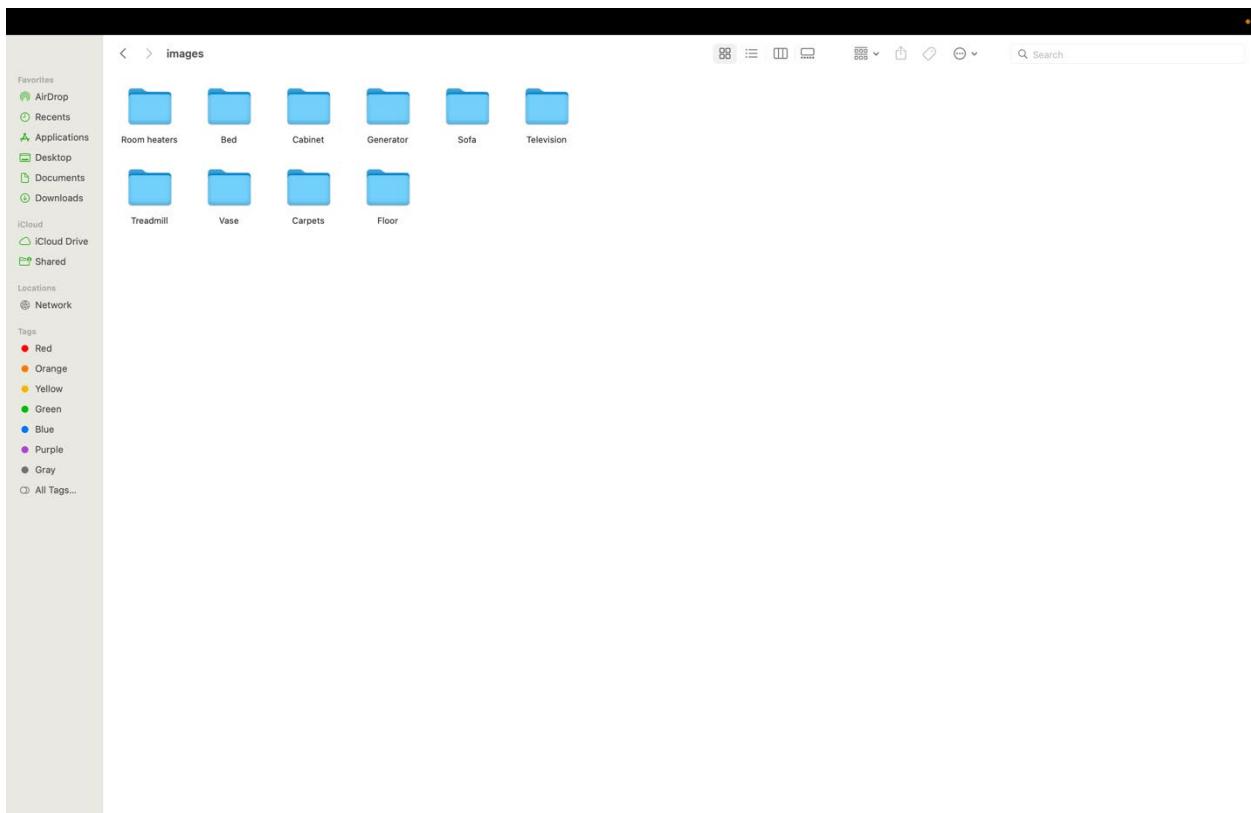
[ ] def fetch_item_images(item_name):
    global api_count
    create_dir(item_name)
    try:
        fetch_item_images_helper(item_name,api_count)
        print(f"Completed for {item_name}")
    except:
        try:
            if api_count < len(api_keys)-1:
                api_count = api_count + 1
                print(f"Changed API to {api_count}")
                fetch_item_images_helper(item_name,api_count)
            else:
                return 0
        except:
            return 0

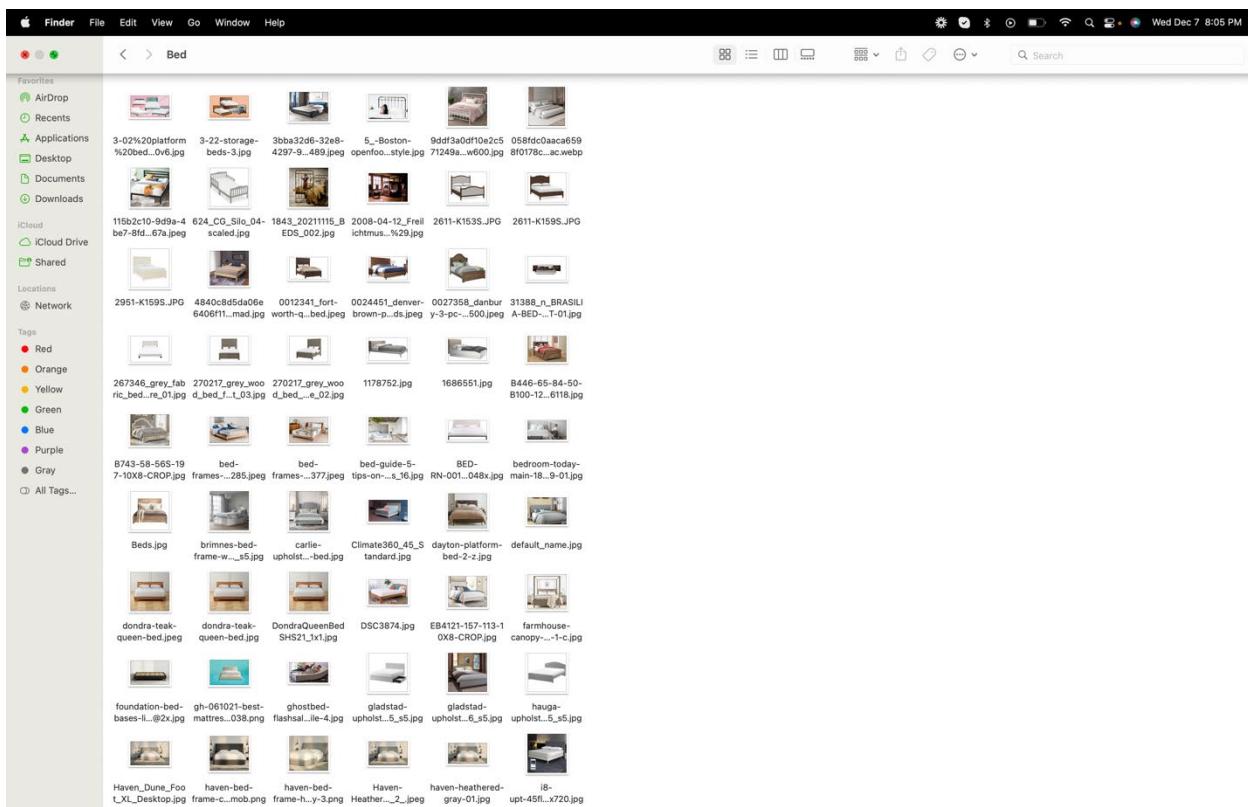
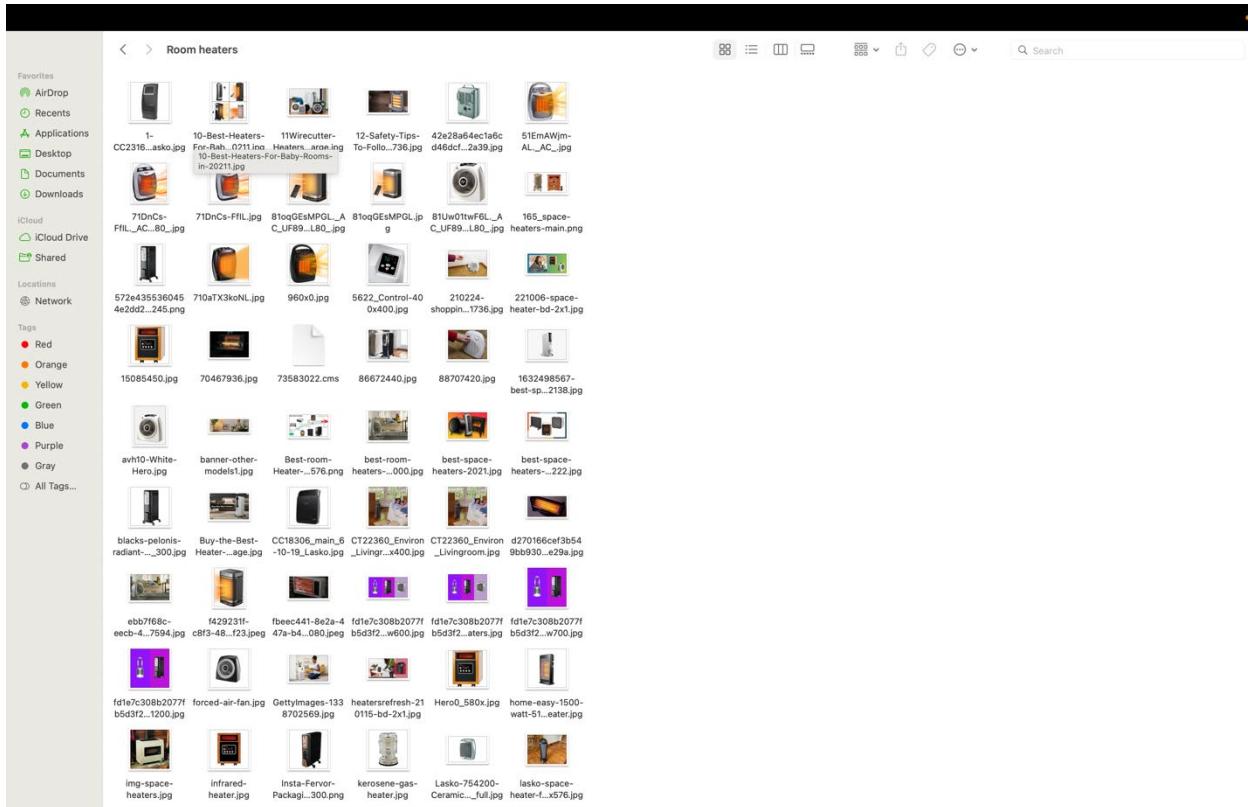
[ ] create_dir("images", item = False)
for each_item in list_of_items:
    fetch_item_images(each_item)

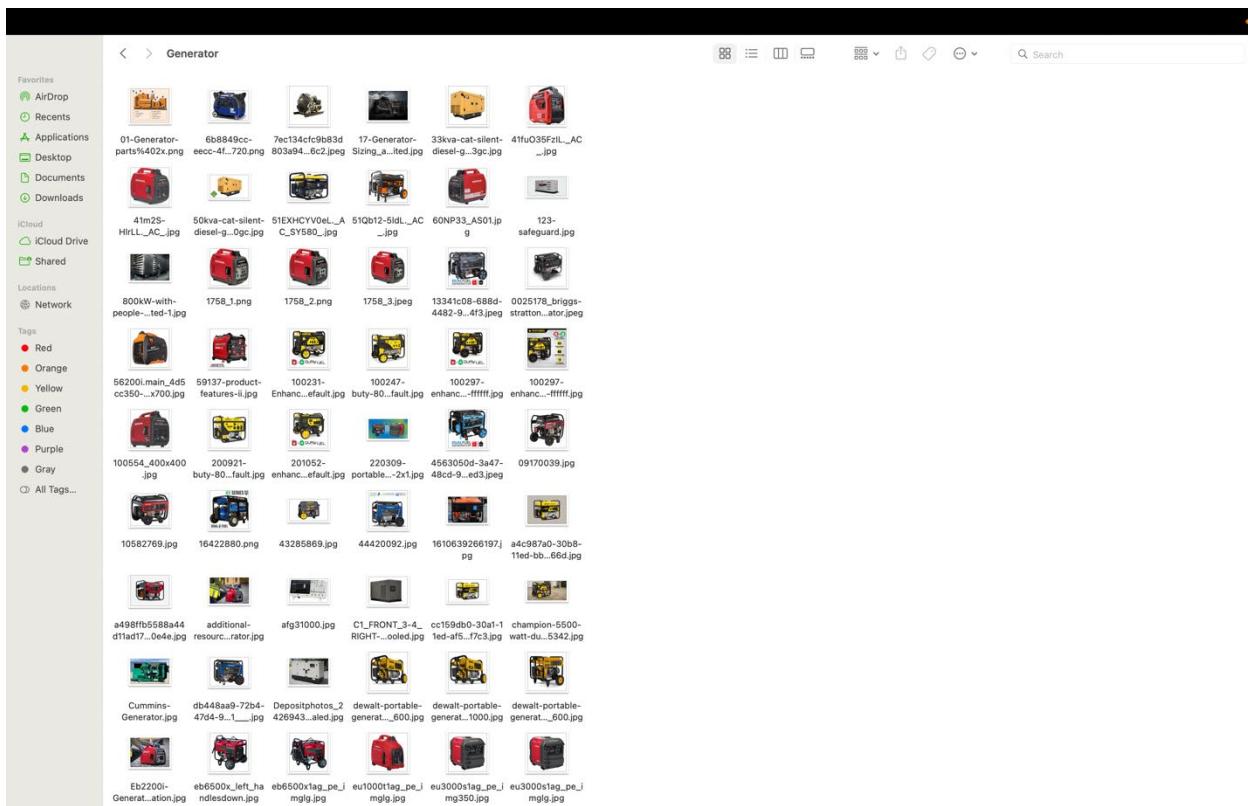
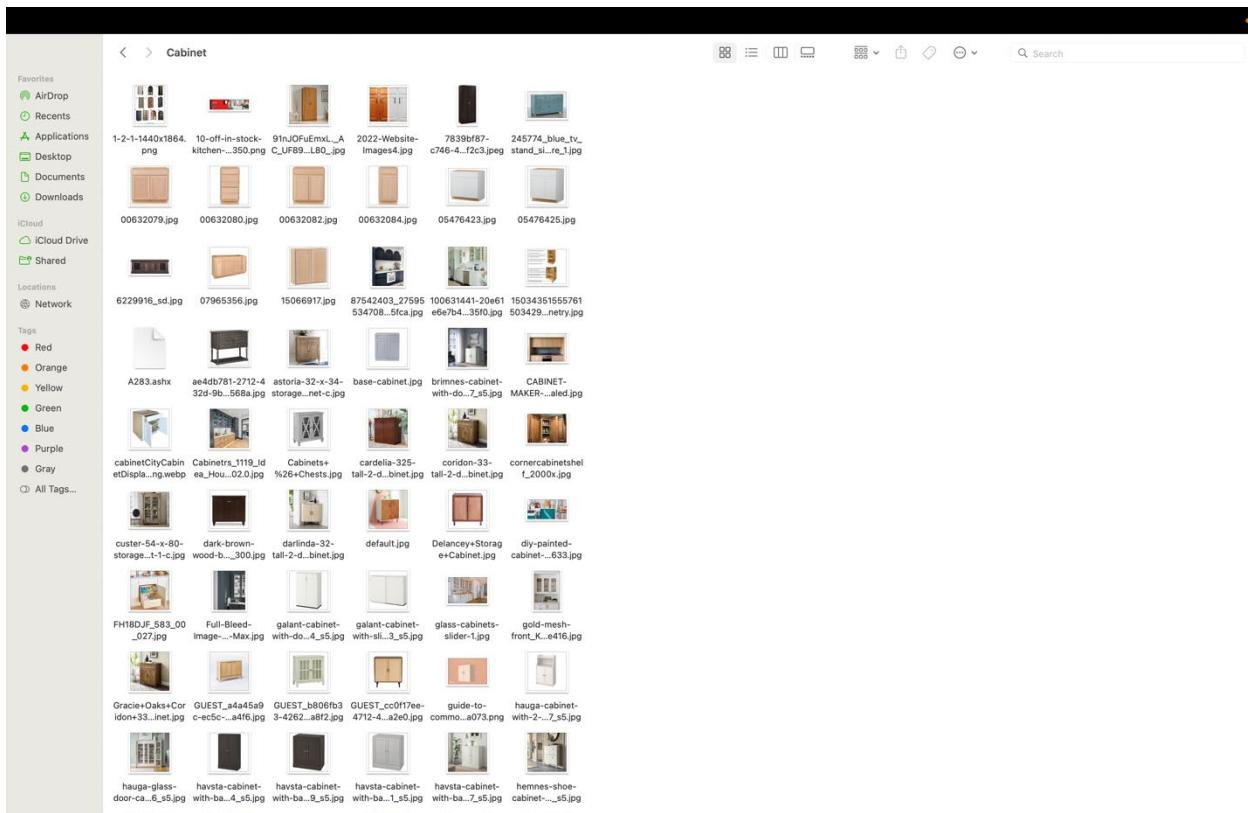
Completed for Room heaters
Completed for Bed
Completed for Television
Completed for Cabinet
Completed for Treadmill
Completed for Vase
Completed for Generator
Completed for Sofa
Completed for Carpets
Completed for Floor
```

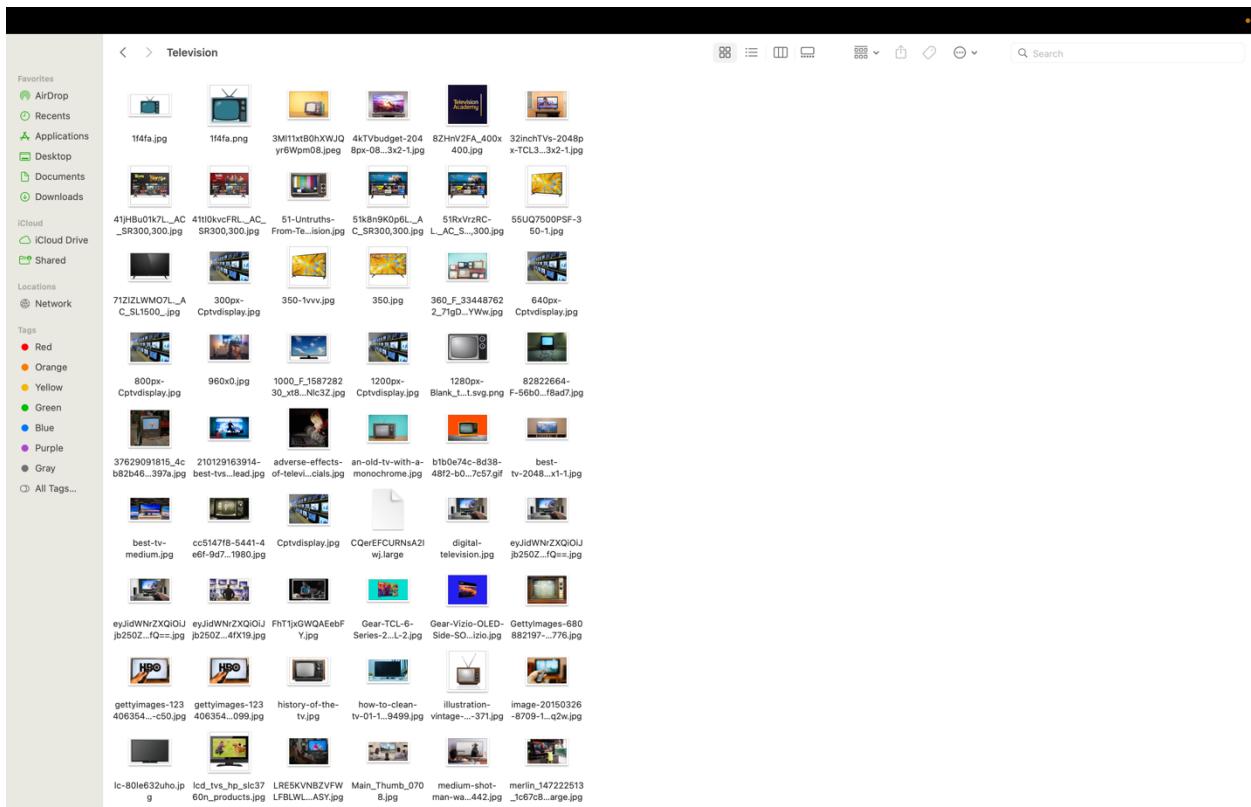
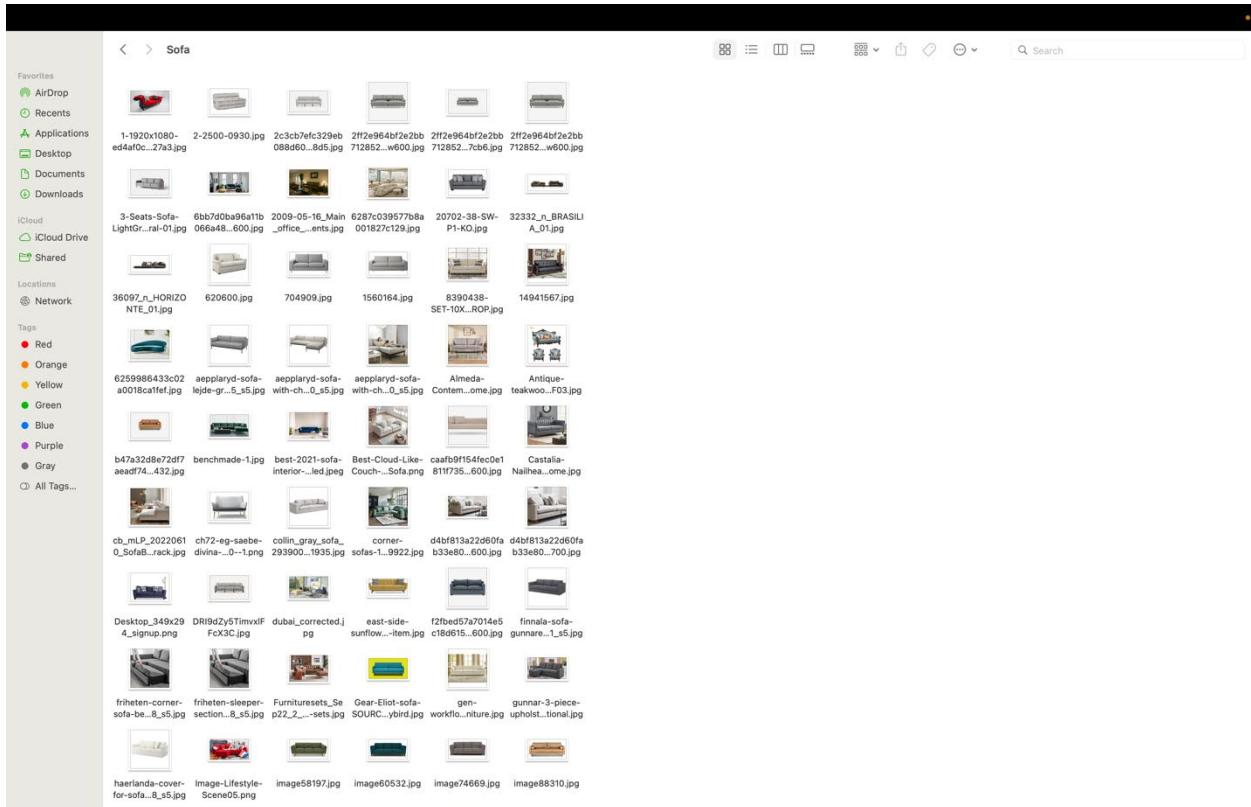
Output











MILESTONE-3

5.3. Milestone 3: Annotation (20 points)

Upload the images in your CVAT instance and annotate all object categories. This is a manual step and you will need to use the [DEXTR cutter](#). Write a 2 page summary on how DEXTR works that can be understood by experts. Ensure that the annotations can be read and seen by the [fiftyone tool](#). Submit a Github branch titled ‘annotation’ containing the annotated dataset in MS COCO format and a markdown file called [annotation.md](#) containing the DEXTR description and the link with the 1000 annotated images (10 product categories x 100 images per category). You can collaborate with other teams on the annotation task to share the load.

<https://universe.roboflow.com/annotations-vzddc/annotations-juxop>

DEXTR is an AI-powered tool that is available to you from the start. It works by you placing four markers (or more) - one on the most northern edge of what you want to label, and one on the most eastern, northern, and western edges too. It is a powerful tool that can get you really good results in just four clicks. DEXTR requires a minimum of four markers to function. When your mouse pointer is over one of these edges, left-click or press "A" to create these markings. The DEXTR-model will "look" for any object inside of those four markers after you have provided four markers. You can keep adding points if you're unsatisfied with the choice. You should be aware, though, that occasionally, the outcomes might not be the best. In our experience, adding more over 10 frequently results in subpar performance, even for complicated objects. By left-clicking, holding, and dragging markers to the correct location, you can move them. You can remove the marker you placed last by pressing "backspace". You can also select a marker that you want to delete by first clicking on it and then pressing "backspace". Tolerance When DEXTR has been activated, you will be able to see a "Threshold" in the top left corner. This threshold can be set from 0-100% by you, the user. It controls how strict the model should be when generating results. The higher the threshold, the smaller and more precise the result will be. Reset the tool If you are unhappy with your selection and want to start over, you can press "esc". Convert to object When you are satisfied with your selection you can turn it into an object by pressing "enter", or by clicking the "convert" button in the tool settings toolbar. DEXTR might take a couple of minutes to activate the first time you annotate as well as when you open up a project where no one has been annotating for the last 30 minutes. Dexter's Offering

AI forecasting for energy trading Dexter offers load and imbalance time-series forecasting throughout the mid-term and short-term trading cycle. Big data management Dexter combines large amounts of different external data sources (i.e. weather data with customer data (historic time-series) to give the most accurate forecast possible. Cloud and API-based services The forecasting engine runs in the cloud and is able to on-board and connect with existing IT infrastructure in a seamlessly easy way.

universe.roboflow.com

NJ Institute of Technology Mail 5. Transfer Learning for Custom... Bathroom Ideas | The Home Dep... dataacquisitionmilestone2.ipynb... DATA-MINING-PROJECT20/DEx... ANNOTATIONS Dataset > Overview Harshitha Saripalli : ANNOTATIONS

Received Links Shared with You Collected Links Bookmarks Reading List

roboflow Projects Universe Documentation Forum

Search 100,000+ Open Source Computer Vision Projects...

★ ANNOTATIONS Computer Vision Project

[Download this Dataset](#)



SOURCE ANNOTATIONS »

LAST UPDATED 15 days ago

PROJECT TYPE Object Detection

SUBJECT BASEMENT

CLASSES Bedrotation, Cabinetrotation, Generatorrotation, Room Heatersrotation, Sofarotation, Televisionrotation, Treadmillrotation, Vaserotation

LICENSE CC BY 4.0 »

Object Detection

A description for this project has not been published yet

Cite this Project

If you use this dataset in a research paper, please cite it using the following BibTeX:

```
@misc{ annotations-juxop_dataset,
    title = { ANNOTATIONS Dataset },
    type = { Open Source Dataset },
    author = { ANNOTATIONS },
    howpublished = { url{ https://universe.roboflow.com/annotations-vzddc/annotations-juxop } },
    journal = { Roboflow Universe },
    publisher = { Roboflow },
    year = { 2022 },
    month = { nov },
    note = { visited on 2022-12-07 },
}
```

SIMILAR PROJECTS More like annotations-vzddc/annotations-juxop »

dataset YunFai_lucar Deteksi Object p11_3 Furniture

universe.roboflow.com

NJ Institute of Technology Mail 5. Transfer Learning for Custom... Bathroom Ideas | The Home Dep... dataacquisitionmilestone2.ipynb... DATA-MINING-PROJECT20/DEx... Project Browser Harshitha Saripalli : ANNOTATIONS

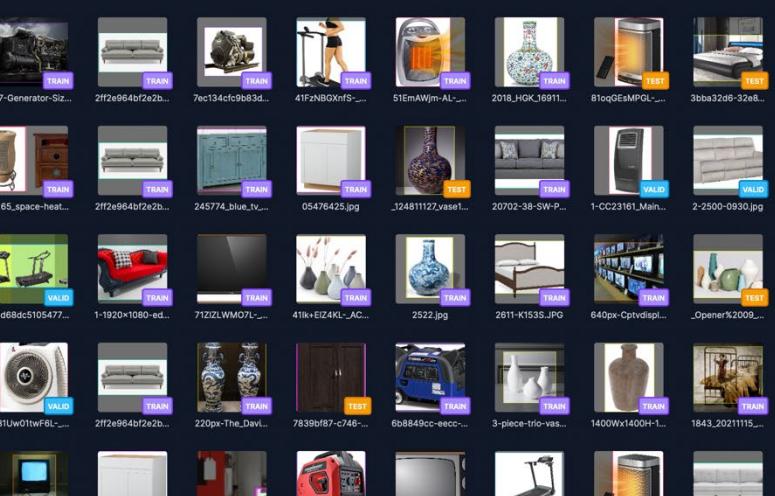
Received Links Shared with You Collected Links Bookmarks Reading List

roboflow Projects Universe Documentation Forum

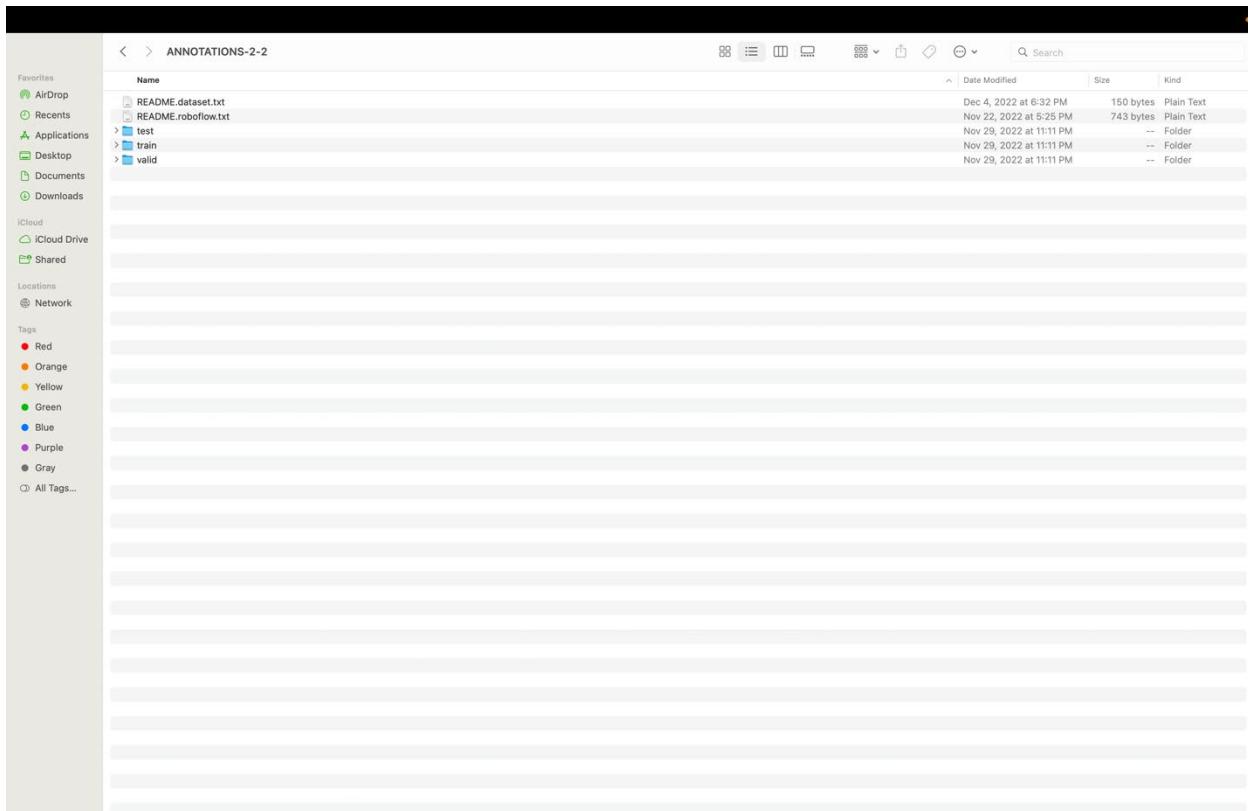
Search 100,000+ Open Source Computer Vision Projects...

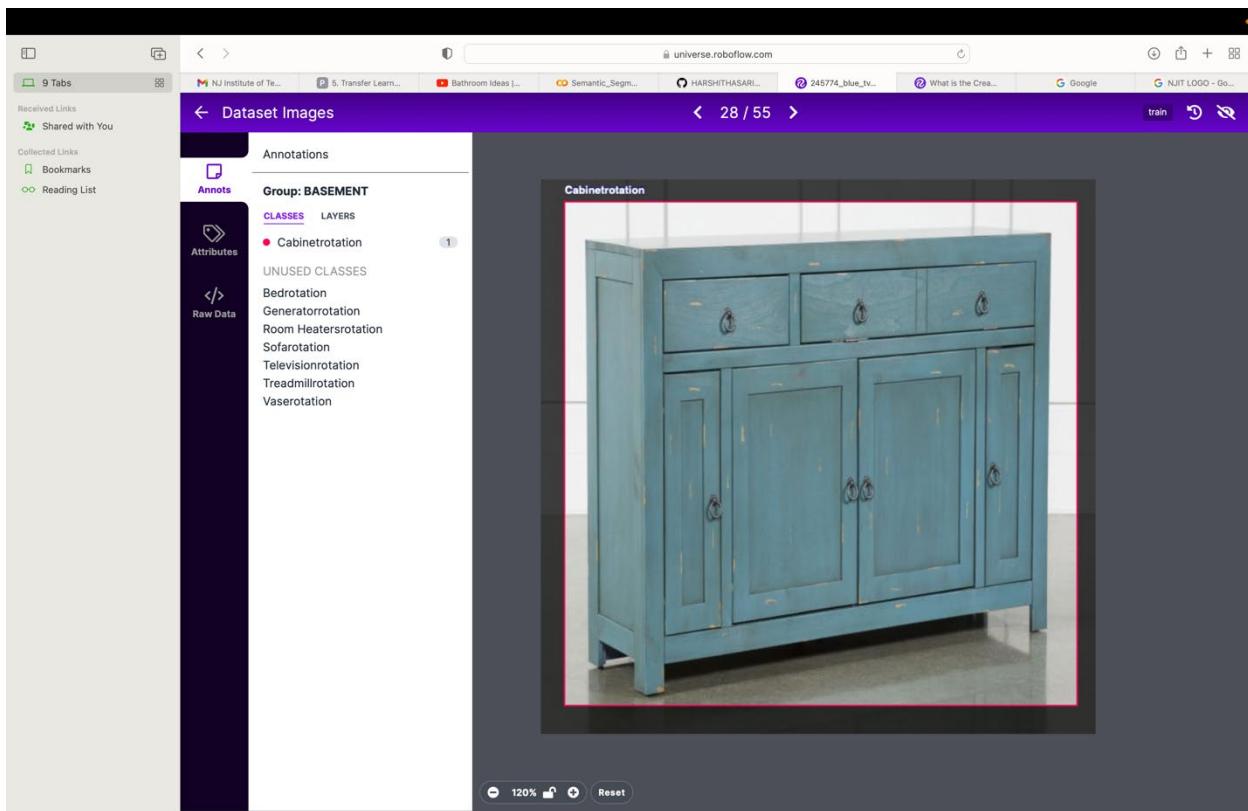
Browse How to Search De-Sel ect Select All Select Images to Clone Search

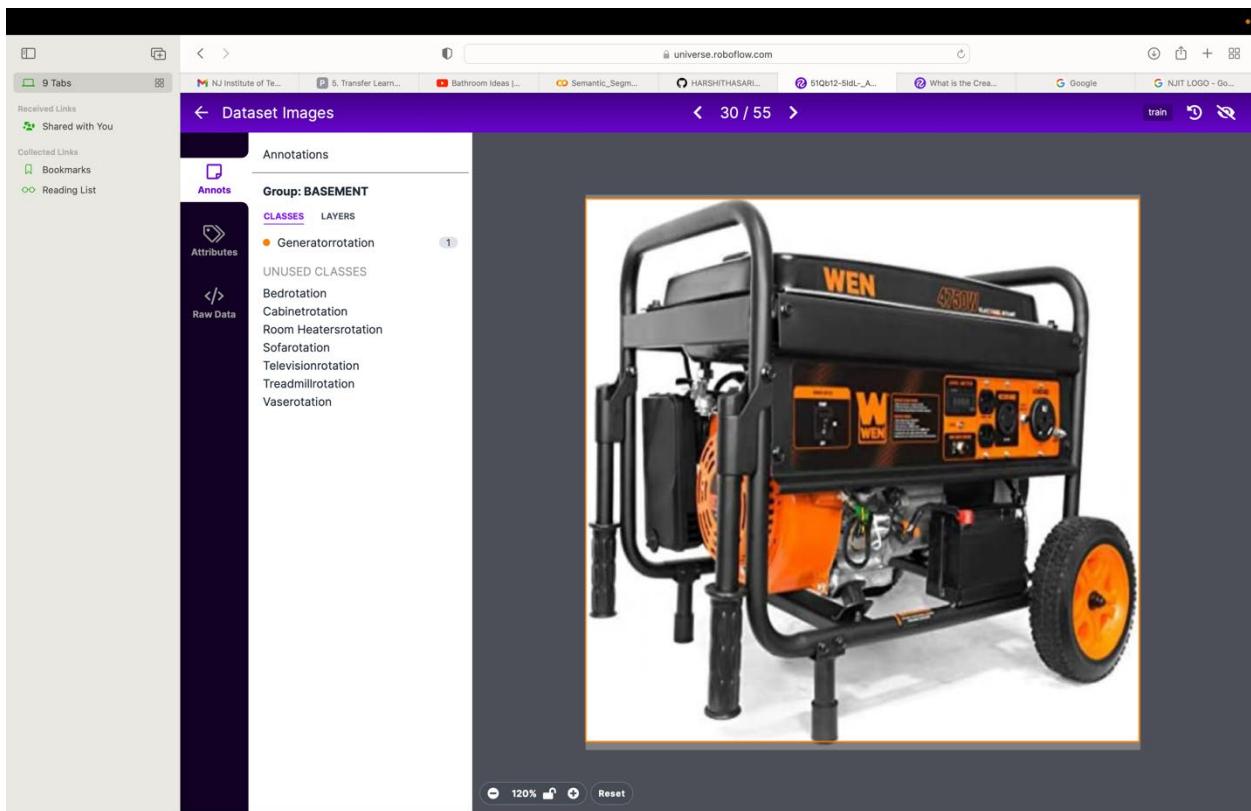
Filename : Split : ANY TRAIN VALID TEST Classes: ALL CLASSES



17-Generator-Siz... 2ff2e964bf2e2b... 7ec134cfcb83d... 41f2NBGXnfS... 51EmAWjm-AL... 2018_HGK_1691... 8ToqGEeMPGL... 3bb32d6-32e...
165_space-heat... 2ff2e964bf2e2b... 245774_blue_tv... 05476425.jpg _24811127_vase1... 20702-38-SW-P... 1-CC23161_Main... 2-2500-0930.jpg
1d68dc510547... 1-1920x1080-ed... 71ZiLWM07L... 41kxEIz4KL-AC... 2522.jpg 2611-K153S.JPG 640px-Optyvdispl... _Opener%2009...
81UwOItwF6L... 2ff2e964bf2e2b... 220px-The_Davi... 7839bf87-c746... 6b8849cc-eecc... 3-piece-trio-vas... 1400Wx1400H-1... 1843_2021115...







Screenshot of the Roboflow web interface showing the "Annotations" dataset page. The "Dataset" tab is selected in the sidebar. A modal window titled "Export" is open, showing options for "Format": COCO, "Annotations used with": EfficientDet Pytorch and Detectron 2, and "download zip to computer" (radio button selected). Other formats shown are MT-YOLOv6, COCO JSON, TFRecord, and CreateML JSON. Below the export modal, there is a preview of 79 images from the dataset, followed by a "View All Images" link. At the bottom, a "TRAIN / TEST SPLIT" section shows the distribution of images: Training Set (55 images, 70%), Validation Set (16 images, 20%), and Testing Set (8 images, 10%).

Screenshot of the Roboflow web interface showing the "Annotations" dataset page. The "Dataset" tab is selected in the sidebar. The main content area displays the dataset details. Under "IMAGES", there is a preview of 79 images and a "View All Images" link. Below this, the "TRAIN / TEST SPLIT" section shows the same distribution as the previous screenshot. Under "PREPROCESSING", it says "Auto-Orient: Applied" and "Resize: Stretch to 640x640". Under "AUGMENTATIONS", it says "No augmentations were applied." At the bottom, the "DETAILS" section provides version information: Version Name: 2022-11-22 12:24pm, Version ID: 1, and Generated: Nov 22, 2022.

MILESTONE-4

5.4. Milestone 4: Semantic Segmentation (40 points)

Use [this](#) template or an equivalent template if you are using Tensorflow to implement a semantic segmentation pipeline based on either MaskRCNN (project team number is odd number) or UNet (project team number is even) that will successfully segment the selected objects on the input room video. You can use frameworks such as <https://github.com/facebookresearch/detectron2> for this task. Ensure you describe how you tuned and what values you ended up using for the [network hyperparameters](#). Submit a Github branch titled 'segmentation' containing the segmentation code and a markdown file called [segmentation.md](#) containing the full description of the semantic segmentation network you used and the performance results you got.

Semantic image segmentation seeks to assign a class to each pixel in an image that corresponds to the concept it is intended to convey. This task is sometimes known as "dense prediction" because we are making predictions for each and every pixel in the image.

The fact that we are just interested in the category of each pixel is crucial to keep in mind because we are not distinguishing between instances of the same class. To put it another way, the segmentation map does not automatically recognize two objects of the same category as independent ones if they are present in the input image. Instance segmentation models are a different class of models that can discriminate between different items belonging to the same class. Each "block" in the architecture is represented by a set of convolution operations in the typical U-Net model. There are a variety of more sophisticated "blocks" that can be used instead of stacked convolutional layers, as I covered in my piece on popular convolutional network topologies.

Drozdzal et al. substitute residual blocks for the fundamental stacked convolution blocks. In addition to the long skip connections that already exist between the relevant feature maps of the encoder and decoder modules in the normal U-Net layout, this residual block also includes short skip connections (inside the block). They claim that training can be completed more quickly and with deeper models because of the brief skip connections.

Jegou et al. elaborated on this by suggesting the use of dense blocks, still having a U-Net structure, saying that the "characteristics of DenseNets make them a very good fit for semantic segmentation as they naturally produce skip connections and multi-scale supervision." These dense blocks are advantageous because they allow for very effective feature reuse by carrying lower level features from earlier layers side by side with higher level features from more recent layers.

Facebook AI Research (FAIR) came up with this advanced library, which gave amazing results on object detection and segmentation problems. Detectron2 is based upon the maskrcnn benchmark. Its implementation is in PyTorch. It requires CUDA due to the heavy computations involved. It supports multiple tasks such as bounding box detection, instance segmentation, keypoint detection, densepose detection, and so on. It provides pre-trained models which you can easily load

Many pre-trained models of Detectron2 can be accessed at [model zoo](#). These models have been trained on different datasets, and are ready to be used.

Even when people are training their custom dataset, they use these pre-trained weights to initialize their model. It has proven to reduce the training time and improve the performance. The model we'll be using is pretrained on the COCO dataset.

First, we have to define the complete configuration of the object detection model. We imported the 'get_cfg' function from the detectron2.config module, we will be using it now. I have chosen the Coco Instance segmentation configuration (YAML file). There are other options available too. You also have to set the model's threshold score (usually set between 0.4 to 0.6). You can load the pretrained weights for the configuration from the checkpoint

The screenshot shows a Google Colab notebook titled "Semantic Segmentation.ipynb". The code cell contains Python code for setting up a Google Drive and cloning a GitHub repository for Detectron2. The output of the cell shows the results of pip installations and dependency resolution.

```
from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

!python -m pip install pyyaml==5.1
import sys, os, distutils.core

!git clone 'https://github.com/facebookresearch/detectron2'
dist = distutils.core.run_setup("./detectron2/setup.py")
!python -m pip install {'.join(["{}" for x in dist.install_requires])}
sys.path.insert(0, os.path.abspath("./detectron2"))

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: pyyaml==5.1 in /usr/local/lib/python3.8/dist-packages (5.1)
fatal: destination path 'detectron2' already exists and is not an empty directory.
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: Pillow>=7.1 in /usr/local/lib/python3.8/dist-packages (7.1.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.8/dist-packages (3.2.2)
Requirement already satisfied: pycocotools>=1.0.2 in /usr/local/lib/python3.8/dist-packages (2.0.6)
Requirement already satisfied: tensorboard in /usr/local/lib/python3.8/dist-packages (1.1)
Requirement already satisfied: tabulate>0.1.8 in /usr/local/lib/python3.8/dist-packages (0.1.8)
Requirement already satisfied: tabulate in /usr/local/lib/python3.8/dist-packages (0.8.10)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.8/dist-packages (1.5.0)
Requirement already satisfied: tqdm>=4.29.0 in /usr/local/lib/python3.8/dist-packages (4.64.1)
Requirement already satisfied: tensorflow in /usr/local/lib/python3.8/dist-packages (2.9.1)
Requirement already satisfied: fvcore<0.1.6,>=0.1.5 in /usr/local/lib/python3.8/dist-package (0.1.5.post20221122)
Requirement already satisfied: iopath<0.1.10,>=0.1.1 in /usr/local/lib/python3.8/dist-packages (0.1.9)
Requirement already satisfied: torch<1.10.0,>=1.9.0 in /usr/local/lib/python3.8/dist-packages (1.9.0)
Requirement already satisfied: pydot in /usr/local/lib/python3.8/dist-packages (1.3.0)
Requirement already satisfied: omegaconf<2.1 in /usr/local/lib/python3.8/dist-packages (2.2.3)
Requirement already satisfied: hydra-core>=1.1 in /usr/local/lib/python3.8/dist-packages (1.2.0)
Requirement already satisfied: black==22.3.0 in /usr/local/lib/python3.8/dist-packages (22.3.0)
Requirement already satisfied: tim in /usr/local/lib/python3.8/dist-packages (0.6.12)
Requirement already satisfied: fairscale in /usr/local/lib/python3.8/dist-packages (0.4.12)
Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages (2.3)
Requirement already satisfied: packaging_ext>=0.10.0 in /usr/local/lib/python3.8/dist-packages (0.10.0)
Requirement already satisfied: platformdirs>=2 in /usr/local/lib/python3.8/dist-packages (from black==22.3.0) (4.1.1)
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.8/dist-packages (from black==22.3.0) (8.1.5)
Requirement already satisfied: mpy-extensions>=0.4.3 in /usr/local/lib/python3.8/dist-packages (from black==22.3.0) (0.4.3)
Requirement already satisfied: tomli>=1.1.0 in /usr/local/lib/python3.8/dist-packages (from black==22.3.0) (2.0.1)
Requirement already satisfied: pathspec>=0.9.0 in /usr/local/lib/python3.8/dist-packages (from black==22.3.0) (0.10.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (from pycocotools==2.0.2) (1.23.5)
Requirement already satisfied: pyparsing<2.0.4,>=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib) (3.0.9)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib) (0.11.0)
Requirement already satisfied: numthon<asturil>=2.1 in /usr/local/lib/python2.7/dist-packages (from matplotlib) (2.0.2)
```

The screenshot shows a Google Colab notebook titled "Semantic Segmentation.ipynb". The code cell contains the following Python script:

```
[ ] import torch, detectron2
!nvcc --version
TORCH_VERSION = ".".join(torch.__version__.split(".")[:2])
CUDA_VERSION = torch._version_.split("+")[-1]
print("torch: ", TORCH_VERSION, "; cuda: ", CUDA_VERSION)
print("detectron2: ", detectron2.__version__)

nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2021 NVIDIA Corporation
Built on Sun_Feb_14_21:12:58_PST_2021
Cuda compilation tools, release 11.2, V11.2.152
Build cuda_11.2.r11.2/compiler.29618528_0
torch: 1.12 ; cuda: cu113
detectron2: 0.6

[ ] import detectron2
from detectron2.utils.logger import setup_logger
setup_logger()

import numpy as np
import os, json, cv2, random
from google.colab.patches import cv2_imshow

from detectron2 import model_zoo
from detectron2.engine import DefaultPredictor
from detectron2.config import get_cfg
from detectron2.utils.visualizer import Visualizer
from detectron2.data import MetadataCatalog, DatasetCatalog

[ ] cfg = get_cfg()
model = "COCO-PanopticSegmentation/panoptic_fpn_R_50_3x.yaml"
cfg.merge_from_file(model_zoo.get_config_file(model))
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5

cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url(model)
predictor = DefaultPredictor(cfg)

[ ] def get_mappings():
    mappings = {}
    mappings[1] = 59
    mappings[7] = 57
    mappings[8] = 62
    mappings[10] = 75
```

Semantic Segmentation.ipynb

```
[ ] import json

with open('/content/drive/MyDrive/Segmentation/ann/annotations/instances_default.json', 'r') as f:
    data = f.read()

ann = json.loads(data)
categories = ann["categories"]
images = ann["images"]
annotations = ann["annotations"]

cat_images_data = {}
img_data = {}

for img in images:
    img_data[img["id"]] = img["file_name"]

for ann_data in annotations:
    cat_id = ann_data["category_id"]
    img_id = ann_data["image_id"]
    if cat_id not in cat_images_data:
        cat_images_data[cat_id] = []
    cat_images_data[cat_id].append(img_data[img_id])

print(categories)

[{'id': 1, 'name': 'Bed', 'supercategory': ''}, {'id': 2, 'name': 'Cabinet', 'supercategory': ''}, {"id": 3, "name": "Carpets", "supercategory": ""}, {"id": 4, "name": "Chair", "supercategory": ""}, {"id": 5, "name": "Dishwasher", "supercategory": ""}, {"id": 6, "name": "Dresser", "supercategory": ""}, {"id": 7, "name": "Fridge", "supercategory": ""}, {"id": 8, "name": "Oven", "supercategory": ""}, {"id": 9, "name": "Shower_curtain", "supercategory": ""}, {"id": 10, "name": "Sofa", "supercategory": ""}]

# for cat in [10]:
#     count = 0
#     print('Cat: ', cat)
#     for img_file_name in cat_images_data[cat]:
#         print(img_file_name)
#         im = cv2.imread("/content/drive/MyDrive/Segmentation/ann/images/" + str(img_file_name))
#         outputs = predictor(im)
#         v = Visualizer(im[:, :, ::-1], MetadataCatalog.get(cfg.DATASETS.TRAIN[0]), scale=1.2)
#         v = v.draw_instance_predictions(outputs["instances"].to("cpu"))
#         out = cv2_imshow(out.get_image()[:, :, ::-1])
#         print(outputs["instances"].pred_classes)
#         print(outputs["instances"].pred_boxes)
#         count += 1
#         if count > 10:
#             break
```

Semantic Segmentation.ipynb

```
[ ] correct = 0
total = 0
for cat in cat_images_data:
    print('Running for category: ', cat)
    for img_file_name in cat_images_data[cat]:
        im = cv2.imread("/content/drive/MyDrive/Segmentation/ann/images/" + str(img_file_name))
        outputs = predictor(im)
        v = Visualizer(im[:, :, ::-1], MetadataCatalog.get(cfg.DATASETS.TRAIN[0]), scale=1.2)
        # v.draw_instance_predictions(outputs["instances"].to("cpu"))
        # cv2.imshow(out.get_image()[:, :, ::-1])
        if cat in get_mappings() and get_mappings()[cat] in outputs["instances"].pred_classes:
            correct += 1
    total += 1

Running for category:  6
Running for category:  10
Running for category:  4
Running for category:  7
Running for category:  3
Running for category:  9
Running for category:  8
Running for category:  2
Running for category:  1
Running for category:  5

[ ] print('Accuracy: ', (correct*100)/total)
Accuracy:  35.476190476190474

[ ] # Visual output

DIR_PATH = '/content/drive/MyDrive/Segmentation/ann/images/'
directory = os.fsencode(DIR_PATH)

selected_files = ['victoria-black-vase.jpg', 'viljestark-vase-clear-glass__0640433_pe699813_s5.jpg', 'vinliden-sofa-hakebo-beige__0852744_pe780233_s5.jpg', 'vint
```

